

# Installation Guide

Meeting notes that never leave your machine. Windows, macOS and Linux — installing, first run, and what to do when something goes wrong.

|              |                                                                                                                                                               |
|--------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Version      | 1.3.1                                                                                                                                                         |
| Updated      | 16 September 2026                                                                                                                                             |
| Download     | <a href="https://github.com/antigravitysoham-eng/vlocalhost-core/releases/latest">https://github.com/antigravitysoham-eng/vlocalhost-core/releases/latest</a> |
| Online guide | <a href="https://vlocal.host/install/guide/">https://vlocal.host/install/guide/</a>                                                                           |
| Support      | <a href="https://vlocal.host/support/">https://vlocal.host/support/</a>                                                                                       |
| Licence      | AGPL-3.0 — free forever, source open                                                                                                                          |

## Contents

|   |                  |                                                                      |
|---|------------------|----------------------------------------------------------------------|
| 1 | Before you start | What the download needs from you.                                    |
| 2 | Windows          | Windows 10 or 11, 64-bit.                                            |
| 3 | macOS            | macOS 12 Monterey or newer.                                          |
| 4 | Linux            | 64-bit, glibc 2.35 or newer (Ubuntu 22.04+, Debian 12+, Fedora 36+). |
| 5 | First run        | Two questions, both changeable later.                                |
| 6 | Troubleshooting  | The failures people actually hit, and what each one means.           |
| 7 | Getting help     | What to send, and where.                                             |

---

# Before you start

What the download needs from you.

Less than you would expect. The installer carries its own Python and every library it needs, so there is nothing to install first.

|            |                                                                                        |
|------------|----------------------------------------------------------------------------------------|
| Disk space | About 400 MB, plus ~150 MB the first time you record, when the speech model downloads. |
| Microphone | Any built-in one is fine.                                                              |
| Internet   | For the download only. Nothing is sent afterwards.                                     |
| Account    | None. There is no sign-in.                                                             |
| Ollama     | Optional — only if you want written summaries as well as transcripts.                  |

*Audio and transcripts never leave your computer. Without Ollama you still get a complete, timestamped transcript of every meeting.*



## Windows

Windows 10 or 11, 64-bit.

### Install

1. Download `vlocalhost-1.3.1-windows-x64-setup.exe` from the downloads page.
2. Double-click it.
3. **Windows will show a blue “Windows protected your PC” screen.** This is expected — see the note below. Click **More info**, then **Run anyway**.
4. Choose where to install it. The default sits inside your user folder and needs no administrator rights. Any drive works — `D:\Apps\Vlocalhost` is fine.
5. Leave **Create a desktop shortcut** ticked if you want one.
6. Click **Install**, then **Finish**. The app opens and asks two setup questions.

**Why the warning appears.** Windows shows that screen for any program without a paid code-signing certificate. It is not a virus warning and says nothing about the file's contents. Vlocalhost is open source — you can read every line of what the installer puts on your machine. A signed build is planned.

*The installer needs no administrator password, so it works on a locked-down work laptop.*

## Or: the portable ZIP

Every release also ships `vlocalhost-1.3.1-windows-x64.zip`: the same build with no installer. Take this one if a virus scanner refuses the `.exe`, if Chrome reports “Virus scan failed”, or if you want nothing written outside a single folder.

1. Unpack the `.zip` anywhere you can write. Keep the path short — a folder buried deeper than about 140 characters cannot unpack cleanly, because the files inside are already long.
2. Open the folder and run `Vlocalhost.cmd`.
3. That is the whole install. Nothing is written outside this folder.

**No desktop icon?** Correct — the ZIP deliberately touches nothing outside its own folder. A helper in the folder does it for you:

```
Create desktop shortcut.cmd
```

*It adds Vlocalhost.AI to your desktop and Start menu, pointing at the copy you unpacked, and finds the real desktop folder even when OneDrive has moved it.*

## Uninstall

**Settings** → **Apps** → **Installed apps** → **Vlocalhost.AI** → **Uninstall**, or use the shortcut in the Start menu folder. Your meeting notes and settings are deliberately kept. To remove those too, delete this folder afterwards:

```
%LOCALAPPDATA%\Vlocalhost
```

The portable ZIP has no uninstaller — that is what portable means. Delete the folder you unpacked, and run `Remove desktop shortcut.cmd` first if you made shortcuts.



## macOS

macOS 12 Monterey or newer.

## Pick the right build

Apple menu → **About This Mac**, and look at the chip:

| Your Mac says           | Download                                      |
|-------------------------|-----------------------------------------------|
| Apple M1 / M2 / M3 / M4 | <code>vlocalhost-1.3.1-macos-arm64.dmg</code> |
| Intel                   | <code>vlocalhost-1.3.1-macos-x64.dmg</code>   |

---

## Install

1. Download the .dmg and double-click it.
2. Drag **Vlocalhost** onto the **Applications** folder in the window that opens.
3. Eject the disk image, then open **Vlocalhost** from Applications or Launchpad.
4. **macOS will refuse the first launch** — see below.
5. The first time you record, macOS asks for **microphone permission**. Allow it, or the app records silence.

### “Apple could not verify Vlocalhost”

macOS blocks apps that have not been notarised by Apple, which requires a paid developer account. Until that is in place, open it manually once:

1. Try to open the app. Dismiss the warning.
2. Open **System Settings** → **Privacy & Security**.
3. Scroll down. There is a line naming Vlocalhost with an **Open Anyway** button. Click it.
4. Confirm. macOS remembers, so you only do this once.

*On older macOS the shortcut was to right-click the app and choose **Open**. That no longer works on macOS 15 and later — use **System Settings**.*

## Recording the other people on a call

macOS has no built-in way to capture what your speakers are playing, so out of the box Vlocalhost records only your microphone. To capture the far end, install **BlackHole** (free, open source), route your output through it, then choose it in **Settings** → **What to listen to**. Windows and Linux need no such step.

## Uninstall

Drag **Vlocalhost** from Applications to the Bin. To remove notes and settings as well:

```
~/Library/Application Support/Vlocalhost
```



## Linux

64-bit, glibc 2.35 or newer (Ubuntu 22.04+, Debian 12+, Fedora 36+).

**Coming soon.** The Linux archive is packaged but not yet published to the releases page. Today, Linux runs from source — see the README in the repository.

Unpack the archive into your home folder and run the install script inside it. It installs for your user only, needs no root, and adds Vlocalhost.AI to the applications menu. The exact commands ship with the build, once that build has been tested on each distribution named above.

The microphone needs **libportaudio2**. The script names the package for your distribution rather than guessing.

Notes and settings are kept in `~/ .local/share/vlocalhost`.

## First run

Two questions, both changeable later.

The first time it opens, Vlocalhost asks two questions. Both can be changed in **Settings** → **Models**, and the whole wizard can be re-run from **Settings** → **Run setup again**.

### 1. How accurate should transcription be?

| Profile  | Model | Memory  | Best for                                                     |
|----------|-------|---------|--------------------------------------------------------------|
| Light    | tiny  | ~250 MB | The default. Leaves the machine free while you are on a call |
| Balanced | base  | ~350 MB | Better transcripts, when nothing else needs the machine      |
| Accurate | small | ~735 MB | 4+ fast cores; noticeably better on Indian languages         |

**Light is the default, and the reason is not accuracy.** A bigger model keeps every core busy for as long as the recording runs, and on some machines that is enough to make **your** voice break up for the other people on a call in a browser — Google Meet, Teams on the web, Zoom in a tab. You cannot hear it; they can. Light leaves enough of the machine free that it does not happen.

Pick **Balanced** or **Accurate** when nothing else needs the machine — recording a room, a call in a desktop app rather than a browser, or transcribing a file afterwards. If people tell you that you sound distorted while you are recording, come back here and choose **Light**.

You can also point it at your own model folder. The model downloads once, the first time you record — expect a wait of a minute or two on that recording only.

### 2. Do you want written summaries?

Summaries are written by a second model running locally through **Ollama**. This is entirely optional:

- **Skip it** and you still get a complete timestamped transcript of every meeting, saved as a .txt file. Nothing else is lost.
- **Set it up** and you additionally get a Markdown file with a summary, key points, decisions and action items.

If Ollama is not installed, get it from **ollama.com**, then return to Settings and press **Run setup again**. The wizard detects it and offers to download the summary model (about 2 GB).

*That is the whole of it for most people. If it does not go that smoothly — or you want to know which model to choose and why — the **Summaries Setup** guide covers Ollama on all three platforms, choosing a model, and every failure this can produce. It ships beside this one, as `vlocalhost-summaries-setup.pdf`.*

## Recording without leaving the meeting

**Ctrl+Shift+Space** starts recording and transcribing, and the same key stops and saves. It works while your call — not Vlocalhost — has focus, which is the whole point: the first half-minute of a meeting is where people say what it is about, and that is the time otherwise spent finding the window.

Change it in **Settings** → **Record without leaving the meeting**. That panel also shows when the key was last pressed, which tells you whether a press reached the app at all — worth knowing on a laptop whose function row is media-locked, where a chord using F1–F12 may never produce a keypress unless you also hold Fn.

*Windows only for now. On macOS and Linux recording starts from the window and the tray as usual; Settings says so rather than leaving you pressing a key that does nothing.*

## Where your files live

|                |                                          |
|----------------|------------------------------------------|
| <b>Windows</b> | %LOCALAPPDATA%\Vlocalhost                |
| <b>macOS</b>   | ~/Library/Application Support/Vlocalhost |
| <b>Linux</b>   | ~/.local/share/vlocalhost                |

These sit deliberately outside the program folder, so updating, reinstalling or uninstalling never touches your meetings. Not sure? The app will tell you:

```
Vlocalhost --paths
```

# Troubleshooting

The failures people actually hit, and what each one means.

## Installing

### Chrome: “Failed — Virus scan failed”

Not a detection. Chrome hands the finished file to Windows to scan, and that call failed.

Usually a third-party antivirus, a managed work laptop, or a Defender service that is not running. Press **Retry** in Downloads first. If it fails again, take the portable ZIP instead — same build, no installer — or verify the file against the SHA-256 published beside every download.

---

## The installer stops part-way and rolls back

The folder you chose is too deep.

Windows limits a path to 260 characters, and the files inside the bundle are already long. An install folder beyond roughly 140 characters cannot fit. Nothing is left behind — install again somewhere shorter, such as the default.

## Windows: antivirus removed the installer

Some scanners flag any unsigned installer that unpacks an interpreter.

Restore it from quarantine and add an exclusion, or verify the download first: each release publishes a SHA-256 checksum, and the .sha256 file next to the download should match what you compute locally.

## macOS: “Vlocalhost is damaged and can't be opened”

Misleading message. It almost always means the quarantine flag is set, not that the file is corrupt.

Use System Settings → Privacy & Security → Open Anyway, as described in the macOS section. If that does not appear, clear the flag directly:

```
xattr -dr com.apple.quarantine /Applications/Vlocalhost.app
```

## macOS: the app closes instantly

Usually the wrong build — an Intel bundle on Apple silicon, or the reverse.

Check the chip under Apple menu → About This Mac and download the matching .dmg.

## Linux: version `GLIBC\_2.35' not found

Your distribution is older than the build.

Install from source instead — the repository README covers it, and it works on much older systems.

## Starting

### Nothing happens when I click the icon

The window failed to open and the error went nowhere, because the shortcut deliberately runs without a console.

Run it from a terminal to see the reason:

```
"%LOCALAPPDATA%\Programs\Vlocalhost\runtime\python.exe"  
"%LOCALAPPDATA%\Programs\Vlocalhost\app\vlocalhost.py"
```

### It opens in a terminal instead of a window

The graphical toolkit could not start, so it fell back.

On Linux this usually means no display is available over SSH. The terminal mode is fully functional — press Ctrl+C to stop and save.

---

## Recording

### The transcript is empty, or every line is silence

The microphone is muted, denied, or the wrong device.

Check the operating system's microphone permission first (macOS: System Settings → Privacy & Security → Microphone; Windows: Settings → Privacy → Microphone). Then confirm the device:

```
Vlocalhost --devices
```

### I can hear the other people but they are not transcribed

A microphone only carries your half of a call.

Set **Settings** → **What to listen to** to **Both**. On macOS this additionally needs BlackHole — see the macOS section.

### The transcript lags behind the meeting

The model is too large for the machine.

Switch to the **Light** profile in Settings → Performance, or press **Benchmark this machine** to measure what your hardware can actually keep up with.

### Hindi (or another language) comes out as nonsense English

An English-only model was selected.

Models ending in .en are English only, and they hallucinate rather than fail. Choose a model without the suffix. The app blocks the combination where it can.

## Summaries

The two that account for nearly all of them are below. The **Summaries Setup** guide has the rest, along with how to install Ollama on each platform.

### The transcript saved but there is no notes file

Ollama was not reachable when the recording stopped.

The transcript is never lost to this — it is written first, on purpose. Check the server is running by opening this in a browser; you should see a short JSON response, not an error:

```
http://localhost:11434/api/tags
```

Then **Settings** → **Models** → **Re-check Ollama**.

### Settings says the model is not installed, but I installed it

Ollama names models with a tag, so llama3.2 and llama3.2:latest are different strings.

Press **List...** next to the summarisation box. It reads the models straight off the server, so you pick a name that certainly exists rather than typing one.

# Getting help

What to send, and where.

Write a diagnostic report and attach it to your message. It lists your settings, versions and recent log lines — **no meeting audio, transcripts or notes**. Read it before sending; it is a plain text file.

```
Vlocalhost --diagnose
```

|                          |                                                                                                                                                               |
|--------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Support</b>           | <a href="https://vlocal.host/support/">https://vlocal.host/support/</a>                                                                                       |
| <b>Downloads</b>         | <a href="https://github.com/antigravitysoham-eng/vlocalhost-core/releases/latest">https://github.com/antigravitysoham-eng/vlocalhost-core/releases/latest</a> |
| <b>Source and issues</b> | <a href="https://github.com/antigravitysoham-eng/vlocalhost-core">https://github.com/antigravitysoham-eng/vlocalhost-core</a>                                 |