

Summaries Setup

Transcripts work the moment you install. Written summaries need one more piece — a language model on your own machine. There are two ways to get one: let the app download it, or run it through Ollama. This is both, how to choose, and what each failure actually means.

Applies to	Vlocalhost 1.2.2, every edition
You need	Ollama, and one model — about 2 GB
Time	About five minutes, once per machine
Updated	01 September 2026
Installation guide	https://vlocal.host/install/guide/
Support	https://vlocal.host/support/

Contents

1	The short way — let the app fetch it	One button. No second program.
2	What this adds	And what already works without it.
3	Step 1 — Install Ollama	One download. Pick your operating system and skip the other two.
4	Step 2 — Get the model	The app will fetch it for you, with a progress bar.
5	Step 3 — Point the app at it	Two fields, both already correct for almost everyone.
6	Choosing a model	Bigger is slower, and not always better.
7	When summaries do not appear	Every failure on this page leaves the transcript intact.

The short way — let the app fetch it

One button. No second program.

If you just want summaries and do not care how, this is the whole procedure. The app can download a small language model and run it itself. Nothing else to install, no account, and once it has downloaded, no internet is needed again.

- 1. Open **Settings** in the app.
- 2. Under **Summaries by**, choose **Built in**.
- 3. Press **Download**. It is about 1 GB and takes a few minutes.
- 4. That is all. Your next meeting gets a summary.

The model is **Qwen2.5 1.5B Instruct**, published under the Apache-2.0 licence — the link beside the Download button opens the terms. It is stored in your own user folder, so an app update does not fetch it again.

*The setup wizard offers the same choice on its **Written summaries** step, so a new installation never has to come here at all. If the wizard finds Ollama already running with a model, it picks that instead rather than downloading a second one.*

Everything after this section is about the other way — **Ollama** — which is worth reading if you already run it, if you want to choose between many models, or if you want a larger model than the built-in one. If you took the button above, you are finished.

What this adds

And what already works without it.

Vlocalhost records and transcribes with no help from anything else. The speech model is the app's own. Summaries are the one feature that needs a second model — either the one the app can download for you, or a general-purpose model you run yourself through Ollama.

Feature	With no notes model	With one
Recording	Works	Works
Live transcript on screen	Works	Works
Saved .txt transcript	Works	Works
Summary, key points, decisions, action items	—	Saved as a .txt file beside the transcript
Meeting named by what was discussed	— named by date and time	Works

This does not change where your data goes. Ollama runs on this machine and the app reaches it on 127.0.0.1. Your audio and your transcript still never leave the computer.

Through Ollama the whole thing is one download of roughly 2 GB and about five minutes of your attention; the built-in model above is about 1 GB and one button. Either way you do it once per machine.

Step 1 — Install Ollama

One download. Pick your operating system and skip the other two.

Ollama is free and open source, from **ollama.com**. It installs a small background service that programs ask to run models — Vlocalhost being one of them. It is not a Vlocalhost component and we do not bundle it, which is exactly why this is a separate step and not something the installer did for you.

Windows

1. Go to **ollama.com/download** and take the Windows installer.
2. Run `OllamaSetup.exe`. It installs into your own user profile and asks for no administrator rights.
3. It starts itself when the install finishes, and again at every login. Look for its icon in the notification area.

*Windows does not reliably put the `ollama` command on your `PATH`. A Command Prompt answering “ollama is not recognized” does **not** mean the install failed — the background service is very probably running perfectly. Ignore the command line and use the app, as in step 2.*

macOS

1. Take the macOS download from **ollama.com/download**.
2. Open the `.dmg` and drag **Ollama** into Applications.
3. Launch it once. It puts an icon in the menu bar and keeps the service running from then on.

Linux

1. Run the installer script from a terminal.
2. It registers a systemd service, so the server comes back after a reboot.

```
curl -fsSL https://ollama.com/install.sh | sh
```

On a distribution without systemd, run `ollama serve` yourself and arrange for it to start at login.

Check that it worked

Open this address in a browser. You want the words **Ollama is running**, not a connection error:

```
http://localhost:11434
```

That is the same check the app makes, and it is the honest one. Whether the `ollama` command exists on your PATH tells you nothing at all about whether the server is up.

Step 2 — Get the model

The app will fetch it for you, with a progress bar.

Ollama is only the engine. It arrives with no models, and until one is pulled it will accept the connection and then refuse every request. This step is not optional, and skipping it is the most common reason summaries never appear.

Through the app

1. Open Vlocalhost.
2. Go to **Settings** → **Run setup again**.
3. Click through to the **Written summaries** step. It looks for Ollama while you watch.
4. When it reports that Ollama is running but has no models, press **Download llama3.2**. The download is about 2 GB and the bar counts the megabytes as they arrive.
5. Press **Next**, then **Finish**. The model is now selected and saved — there is nothing further to configure.

Start this on a connection you are happy to spend 2 GB on. An interrupted download resumes: begin it again and it continues from where it stopped rather than starting over.

Or from a terminal

Equivalent, and quicker if you already know which model you want:

```
ollama pull llama3.2
```

On Windows, if that command is not recognised, call it by its full path instead:

```
%LOCALAPPDATA%\Programs\Ollama\ollama.exe pull llama3.2
```

Then go back to the app and press **Settings** → **Models** → **Re-check Ollama**, so it notices what you just did.

Pull the model before your next meeting, not during it. Ollama downloads and answers on the same server, so a pull still running when a meeting ends will make that summary crawl.

Step 3 — Point the app at it

Two fields, both already correct for almost everyone.

Open **Settings** → **Models**. Exactly two things there concern summaries:

Field	What it is	Change it only if
Ollama server	Where the server is. Default <code>http://localhost:11434</code>	You moved Ollama to another port, or you are deliberately using another machine
Summarization	Which pulled model writes the notes	You want a different one — press List... rather than typing

Always choose the model with the **List...** button instead of typing its name. It reads the names straight off the server, so you cannot pick one that does not exist — which is the single most common way to end up with no summaries and no explanation of why.

The line beside the box is the verdict, and it is specific. ✓ **llama3.2 ready** means everything is in place. Press **Re-check Ollama** after any change.

Prove it with a two-minute test meeting: record, say something, press Stop. You should get **two** files in your notes folder, not one — a `-transcript.txt` and a `-summary.txt`.

If you point **Ollama server** at another computer, your transcript is sent to that computer. That may be precisely what you want on your own network — but it is then no longer true that nothing leaves this machine, so make that choice deliberately.

Choosing a model

Bigger is slower, and not always better.

Any model Ollama can run will work. These are the ones worth considering for meeting notes, with what each wants from your machine. Sizes are approximate and shift as the publishers update them.

Model	Disk	Free RAM	Best for
llama3.2:1b	~1.3 GB	~2 GB	An old or small machine. Summaries are serviceable; action items are the first thing to suffer
llama3.2	~2 GB	~4 GB	The default, and what the app offers to fetch. The right first choice on almost any machine
qwen2.5:7b	~4.7 GB	~8 GB	Better structure, and better on non-English meetings, if you have the memory to spare

mistral

~4.1 GB

~8 GB

A solid 7B alternative with the same appetite

What actually decides the wait

Summarizing a real meeting on a local model is genuinely slow — around two minutes for a 7,000-character transcript on llama3.2, on an ordinary laptop. Two things move that number:

- **Model size.** A 7B model is roughly twice the wait of a 3B one on the same hardware.
- **Meeting length.** The transcript is the input, so an hour costs considerably more than ten minutes.

The transcript is written to disk **before** summarizing begins. However long the model takes, and even if it fails outright, the record of the meeting is already safe on disk.

When summaries do not appear

Every failure on this page leaves the transcript intact.

The transcript saved but there is no notes file

Ollama was not reachable when the recording stopped.

Open `http://localhost:11434` in a browser. Nothing there means the server is not running: start Ollama — from the Start menu on Windows, from Applications on macOS, or with `systemctl start ollama` on Linux. Then press **Settings** → **Models** → **Re-check Ollama**.

“Running, but no models pulled”

Ollama is installed and answering, but empty.

This is the normal state of a fresh install, not a fault. Go back to step 2 and pull a model.

Settings says the model is not installed, but I installed it

Ollama names models with a tag, so `llama3.2` and `llama3.2:latest` are different strings.

Press **List...** beside the summarization box and choose from the list instead of typing the name.

The summary takes longer than the meeting did

The model is too large for the machine, or something else is competing for the same memory.

Drop to a smaller model — `llama3.2:1b` is the fastest listed above. Close anything else that wants several gigabytes. A model pull running in the background does this too; let it finish first.

The notes come back empty, truncated, or as nonsense

The model ran out of memory part-way and was cut short.

Choose a smaller model. A machine that cannot hold the whole model in memory pages it in and out, and the answer degrades well before it fails outright.

It worked yesterday and today it does not

Ollama is no longer starting with the machine.

The installers arrange this on Windows and macOS, but a startup-cleanup tool or a workplace policy can undo it. Start Ollama and confirm its icon is in the notification area or menu bar. On Linux: `systemctl enable --now ollama`.

A work laptop blocks the download

The pull goes out to the internet, through whatever proxy your network insists on.

Ollama honours the usual `HTTPS_PROXY` environment variable. Failing that, pull the model on a machine that can reach the internet and copy the Ollama models folder across — the app itself needs no network at any point.

None of this explains it

The diagnostic report names which check failed.

It lists your settings, versions and recent log lines — **no meeting audio, transcripts or notes**. It is a plain text file; read it before you send it anywhere.

```
vlocalhost --diagnose
```

Still stuck

Support	https://vlocal.host/support/
Full installation guide	https://vlocal.host/install/guide/
Ollama documentation	https://ollama.com